

Full Name: James Le

**CS-281, Introduction to Systems**  
**Fall 2015**  
**Take Home Final Exam**

Due by Thursday, December 17, 2015, 3 p.m.

**Instructions:**

- Make sure that your exam is not missing any pages, then write your full name on the front.
- Write your answers in the space provided below the problem. If you make a mess, clearly indicate your final answer.
- Self-imposed *strict* time limit of 3 contiguous hours. You should not read any of the questions until you are ready to begin and you should stop after three hours of work.
- Open textbook, and open professor-provided handouts (Y86 Instruction set and Instruction semantics).
- Calculators may be used, but **only** for addition, subtraction, multiplication or division of decimal numbers.
- **You may not** use any other electronic aids in the completion of this test.
- The problems are of varying difficulty. The point value of each problem is indicated. If you feel any problem has multiple reasonable interpretations, clearly state your assumptions and solve it under those assumptions.
- After reading these instructions, affix your signature to this front page below to indicate your understanding and compliance with the above instructions and the academic integrity statement that follows.

*On my honor, I acknowledge the above instructions and assert that I will follow these instructions and limit myself to three contiguous hours, to not speak with other students about any questions, and to otherwise act with honor and integrity in the completion of this exam.*




Khanh Nom "James" Le

### Problem 1. (15 points):

Suppose, as a digital logic hardware designer, you are given the table below, which describes state transitions and outputs for a desired sequential circuit. In the table,  $Q$  represents the current state, the next state ( $Q'$ ) is given by the next two columns based on the value of binary input  $X$ , and the output  $Z$  associated with each current state is given in the final column. The four states of the system are represented symbolically as  $A$ ,  $B$ ,  $C$ , and  $D$ . Your goal is to design a sequential circuit for this system using 2  $D$  flip-flops to store the current state. Note that, since the output  $Z$  is determined solely by the current state, this state transition table describes a Moore Machine.

$Q$	$X = 0$	$X = 1$	$Z$
A	A	C	1
B	B	D	0
C	C	B	1
D	D	C	0

- A. Begin by drawing the Finite State Machine (Moore style) that corresponds to the described sequential circuit.



- B. Assume that the state assignment for our two D flip flops ( $D_1 D_0$ ) is  $A = 00$ ,  $B = 01$ ,  $C = 10$ ,  $D = 11$ . Fill in the truth tables below for both the output and the "next state" combinational logic:

$D_1$	$D_0$	$Z$
0	0	1
0	1	0
1	0	1
1	1	0


	$D_1$		
$D_0$	0	1	
0	1	1	$\rightarrow Z = \overline{D_0}$
1	0	0	

$D_1$	$D_0$	$X$	$D_1'$	$D_0'$
0	0	0	0	0
0	0	1	1	0
0	1	0	0	1
0	1	1	1	1
1	0	0	1	0
1	0	1	0	1
1	1	0	1	1
1	1	1	1	0

		$D_0 X$			
$D_0'$	$D_1$	00	01	10	11
	0	0	0	1	1
	1	0	1	1	0

$$\Rightarrow D_0' = \bar{X} D_0 + X$$

- C. Using K-maps (and circling necessary prime implicants), design a minimal circuit for only next state variable  $D_0'$  and output  $Z$ , giving the minimal boolean expression in the resultant sum-of-products form. Show your work to the right of the truth tables above. Just give Boolean algebra expressions for  $D_0'$  and for  $Z$ , you **do not** have to draw the gate diagram.



## Problem 2. (12 points):

Assume you have an 8 bit floating point format with a sign bit ( $s = 1$ ),  $k = 4$ , and  $n = 3$ . Fill in the required values below based on the following instructions for each field. Be careful to give me  $M$  and  $E$ , not exp and frac.

- **Binary:** The 8 bit binary representation.
- **M:** The value of the significand. This should be a number of the form  $x$  or  $\frac{x}{y}$ , where  $x$  is an integer, and  $y$  is an integral power of 2. Examples include  $0$ ,  $\frac{5}{4}$ .
- **E:** The base 10 integer value of the exponent.
- **Value:** The base 10 numeric value represented.

Note: you need not fill in entries marked with "—".

Description	Binary	$M$	$E$	Value
—	0 1 0 0 0 1 0 1	$13/8$	$11$	$13/4$
Smallest positive normalized number	1 0 0 0 1 0 0 0	$1$	$-6$	$1/64$
Largest positive normalized number	1 1 1 1 0 1 1 1	$17/8$	$7$	—
Largest positive denormalized number	0 1 0 0 0 1 1 1	$7/8$	$-6$	—
—	1 1 0 1 1 1 0 1 1	$11/8$	$2$	$\frac{11}{2}$

$$\text{Bias} = 2^{k-1} - 1 = 2^3 - 1 = 7$$

Normalized (exp ∈ {1, ..., 14})

Smallest  $E$ :  $-6$  (exp = 1)

Largest  $E$ :  $7$  (exp = 14)

Smallest  $M$ :  $1.000$

Largest  $M$ :  $1.111$

Smallest nor.:  $1.0 \times 2^{-6} = 0.015625$

Largest nor.:  $1.111 \times 2^7$

$$\text{exp} \in \{0, \dots, 15\}$$

Denormalized

• Exp = 0

•  $E = 1 - \text{Bias} = 1 - 7 = -6$

• Largest  $M$ :  $0.111$

• Largest Denormalized:

$$0.111 \times 2^{-6}$$

The following problem concerns the following, low-quality code:

```
void foo(int x)
{
    int a[3]; ✓
    char buf[4]; ✓
    a[0] = 0xF0F1F2F3; ✓
    a[1] = x; ✓
    *gets(buf);
    *printf("a[0] = 0x%x, a[1] = 0x%x, buf = %s\n", a[0], a[1], buf);
}
```

In a program containing this code, procedure `foo` has the following disassembled form on an IA32 machine:

foo:

```
✓ | pushl %ebp save %ebp on stack
  | movl %esp, %ebp
  | pushl %ebx save %ebx
  | subl $32, %esp allocate 32 bytes on stack

  | movl 8(%ebp), %eax retrieve canary
  | ✓ movl $0xf0f1f2f3, -16(%ebp) store on stack // a[0]
  | ✓ movl %eax, -12(%ebp) store on stack

  | leal -20(%ebp), %ebx compute buf as %ebp-20
  | movl %ebx, (%esp) store buf at top of stack
  | (call) 80484a4 <_init+0x54> # gets call "gets"

  | movl %ebx, 12(%esp) store buf at top of stack
  | ✓ movl -12(%ebp), %eax retrieve canary
  | movl %eax, 8(%esp)
  | ✓ movl -16(%ebp), %eax retrieve canary
  | movl %eax, 4(%esp)
  | movl $0x80490ec, 0(%esp) # address of format string
  | (call) 8048494 <_init+0x44> # printf call "printf"

  | ✓ | addl $32, %esp deallocate stack space
  | | popl %ebx restore %ebx
  | | movl %ebp, %esp
  | | popl %ebp restore %ebp
  | | ret return
```

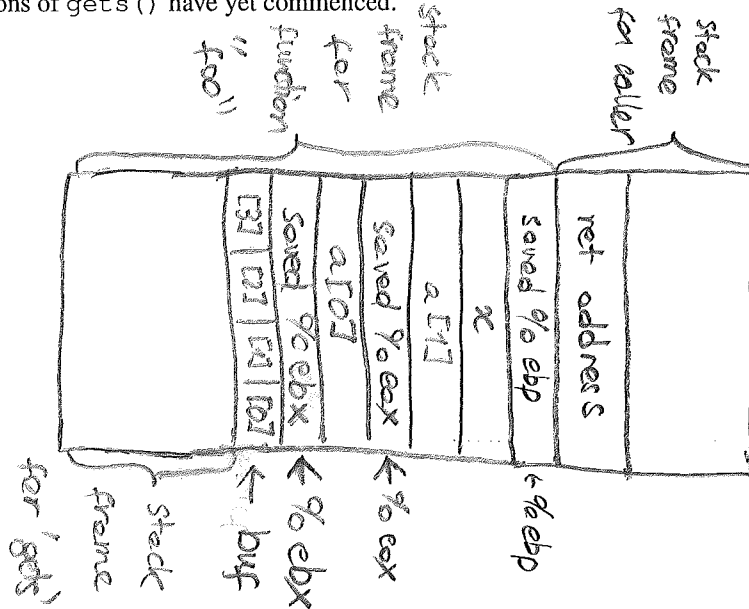
For the following questions, recall that:

- gets is a standard C library routine, like we saw in the buffer lab, with no bounds checking.
- IA32 machines are little-endian.
- C strings are null-terminated (i.e., terminated by a character with value 0x00).
- Characters '0' through '9' have ASCII codes 0x30 through 0x39.

### Problem 3. (12 points):

Consider the case where procedure `foo` is called with argument `x` equal to `0xE3E2E1E0`, and we type `"123456789"` in response to `gets`.

- A. Draw a detailed picture of the stack (using a 4-byte wide representation) with a frame for `foo` labeled with any known locations for variables and/or registers and values, up to the point where `gets()` has just been called, but no instructions of `gets()` have yet commenced.



- B. Fill in the following table indicating which program values are/are not corrupted by the response from `gets`, i.e., their values were altered by some action within the call to `gets`.

Program Value	Corrupted? (Y/N)
<code>a[0]</code>	<del>N</del> Y
<code>a[1]</code>	N
<code>a[2]</code>	N
<code>x</code>	N
Saved value of register <code>%ebp</code>	N
Saved value of register <code>%ebx</code>	Y

- C. What will the `printf` function print for the following:

- `a[0]` (hexadecimal): 0x f0 f1 f2 39
- `a[1]` (hexadecimal): 0x e3 e2 e1 e0
- `buf` (ASCII): 34 33 32 31

$$S = \frac{C}{B \times E} \quad t = m - (s + b) \quad s = \log_2(S) \quad b = \log_2(B)$$


#### Problem 4. (6 points):

The following table gives the parameters for a number of different caches, where  $m$  is the number of physical address bits,  $C$  is the cache size (number of data bytes),  $B$  is the block size in bytes, and  $E$  is the number of lines per set. For each cache, determine the number of cache sets ( $S$ ), tag bits ( $t$ ), set index bits ( $s$ ), and block offset bits ( $b$ ).

Cache	$m$	$C$	$B$	$E$	$S$	$t$	$s$	$b$
1.	32	1024	4	4	64	24	6	2
2.	32	1024	4	256	1	30	0	2
3.	32	1024	8	128	1	29	0	3

#### Problem 5. (6 points):

In our Y86 programs, we encounter many cases in which we want to add a constant value to a register. This requires two instructions, an `irmovl` followed by an `OP1` where the `ifun` indicates addition. Suppose we want to add a new instruction `iaddl` with the following format:



This instruction adds the constant value  $V$  to register  $rB$ . Using Register Transfer Language, specify the six stages, from Fetch to PC Update, to implement this instruction.



⇒ Datapath by Stages for "iaddl"

Fetch	icode: ifun $\leftarrow M_1[PC]$	Execute	val E $\leftarrow$ valC + valB + valA CC: <del>ZF</del> ZF $\leftarrow$ 0, SF $\leftarrow$ 0, OF $\leftarrow$ 0					
	rA: rB $\leftarrow M_1[PC + 1]$	Memory	N/A					
	valC $\leftarrow M_4[PC + 2]$	Write Back	R[rB] $\leftarrow$ val E					
	valP $\leftarrow PC + 8$	PC Update	PC $\leftarrow$ valP					
Decode	valA $\leftarrow$ R[rA]	oddl V, rB <table border="1"><tr><td>C</td><td>0</td><td>F</td><td>rB</td><td><math>\leftarrow</math> V <math>\rightarrow</math></td></tr></table>		C	0	F	rB	$\leftarrow$ V $\rightarrow$
	C			0	F	rB	$\leftarrow$ V $\rightarrow$	
valB $\leftarrow$ R[rB]								

## Problem 6. (14 points):

Consider the function `baz` disassembled below. Assume that `baz` is invoked from a `main` function in a manner consistent with the C code:

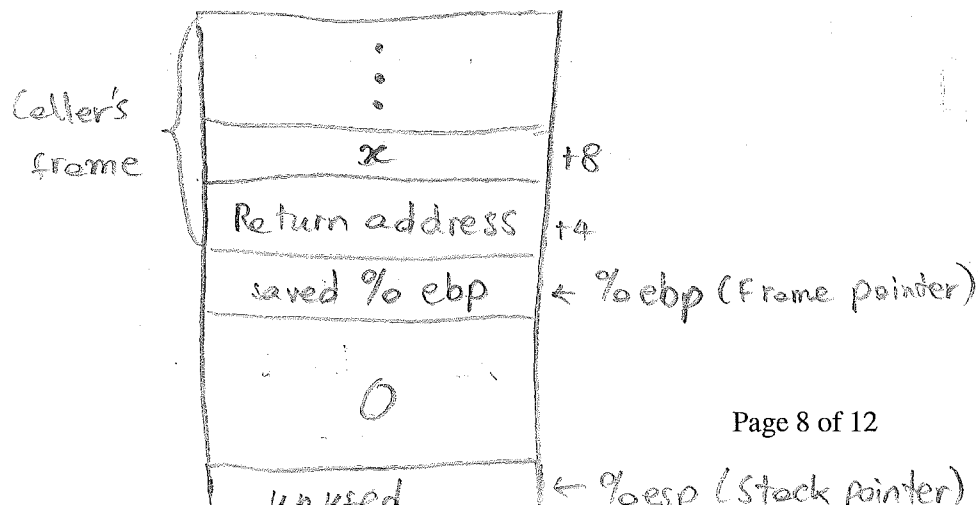
```
int main() {
    int x = 5; ✓
    int y = baz(5); ✓
    printf("baz(%d) = %d\n", x, y); ✓
    exit(0);
}
```

Dump of assembler code for function `baz`:

```
=> 0x0804844f <+0>: push    %ebp    save old %ebp
0x08048450 <+1>: mov     %esp, %ebp    set %ebp as frame pointer
0x08048452 <+3>: sub     $0x8, %esp    allocate 8 bytes on stack
0x08048455 <+6>: cmpl    $0x0, 0x8(%ebp)    compare 0: (jump if > 0)    main's (0+ %ebp)
0x08048459 <+10>: jg      0x8048462 <baz+19>
0x0804845b <+12>: mov     $0x0, %eax    get x=0 (jump anyway)
0x08048460 <+17>: jmp     0x8048479 <baz+42>
0x08048462 <+19>: mov     0x8(%ebp), %eax    get x
0x08048465 <+22>: sub     $0x2, %eax    x-2 (x=2)
0x08048468 <+25>: mov     %eax, (%esp)    store x on stack (on esp)
0x0804846b <+28>: call    0x0804844f <baz>    call "baz"
0x08048470 <+33>: mov     %eax, -0x4(%ebp)    store x on stack (4 down of ebp)
0x08048473 <+36>: mov     0x8(%ebp), %eax    get x
0x08048476 <+39>: add     -0x4(%ebp), %eax    return value = x + (x-2)
0x08048479 <+42>: add     $0x8, %esp    deallocate 8 bytes from stack
0x0804847c <+45>: pop     %ebp    restore %ebp
0x0804847d <+46>: ret     return
```

Your job is to draw, in detail, all stack frames corresponding to the execution of this program, beginning with the stack frame for `main` at the time of the call to `baz`. Also indicate the result of the `printf` in `main` at the conclusion of the program. Do not worry about indicating the completion of stack frames as functions complete. Show the limits of each stack frame and, within the boxes in the stack, indicate the words used in the stack frame. If there is a known value, enter that. If the value is an address, give a short qualitative description (e.g. "ret addr to main", or "main's `%ebp`").

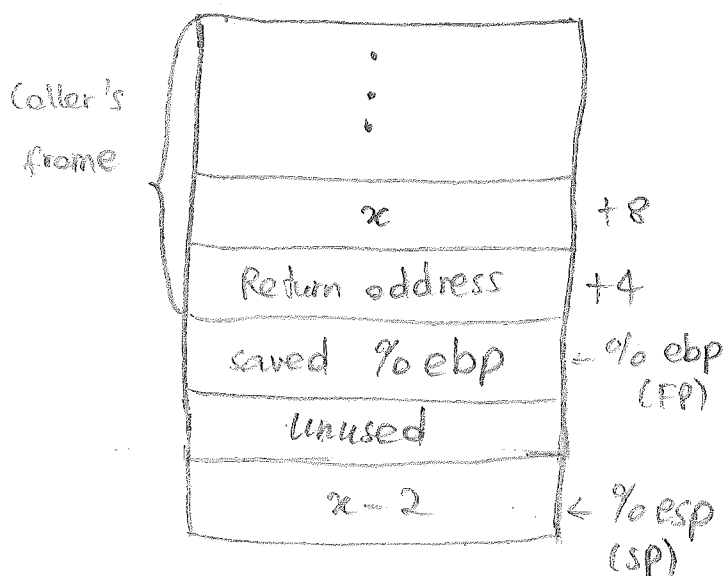
Case I - The argument ~~for baz~~<sup>x</sup> is less than 0. There will be no call to `baz`, and there is only 1 stack frame of the `main` function,



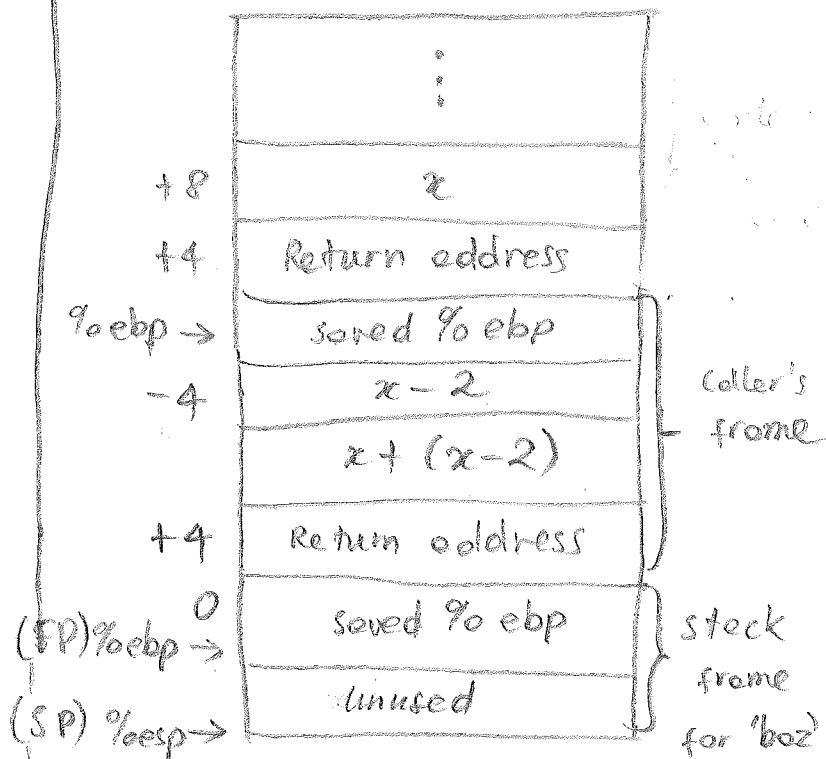


Case II - The argument  $x$  is bigger than 0. There are 2 displayed frames:

\* Frame 1 - Just before call to 'boz'



\* Frame 2 - In body of 'boz'



Main function will print out:

(5, 8)

to the screen

### Problem 7. (6 points):

Consider the following C program:

```
#include <sys/wait.h>

main() {
    int status;
    int pid;

    printf("%s\n", "Hello"); ✓
    pid = fork();
    printf("%s\n", (pid != 0) ? "parent" : "child"); // !fork()

    if (wait(&status) != -1)
        printf("%d\n", WEXITSTATUS(status));

    printf("%s\n", "Bye"); ✓
    exit(2);
}
```

Recall the following:

- Function `fork` returns 0 to the child process and the child's process Id to the parent.
- Function `wait` returns -1 when there is an error, e.g., when the executing process has no child.
- Macro `WEXITSTATUS` extracts the exit status of the terminating process.

Give two valid output results of this program. *Hint: there are several correct solutions.*

Output 1: Hello 0 1 Bye 2 Bye

Output 2: Hello 1 Bye 0 2 Bye

### Problem 8. (15 points):

Consider a memory system with a main memory where  $w = 8$  so there are  $2^8 = 256$  bytes of main memory. Further assume that the memory contents have the values as given in the table on the next page, wherein the left hex digit of an address gives the row of the table and the right hex digit of an address gives the column.

We introduce a level-1 cache where there are four bytes per cache block, the cache has 4 sets, and there are two lines per set (aka, a 2-way set associative cache), and the cache uses an LRU replacement policy on evictions. The cache has initial contents as given below. Your task is to simulate operation of the cache by translating the sequence of address references below, noting *Hit* or *Miss* below the address reference, and updating the cache picture with new values for any updated cache line by writing the new values below any previous cache line values in the cache contents picture below. Assume all address references are *load* operations and access a single, so you do not need to worry about write policies, nor an access crossing a cache line boundary.

$$E = 2$$

$$B = 4 \rightarrow b = 2$$

$$S = 4 \rightarrow s = 2$$

$$t = 4$$

C7	C7	C7	C7	CI	CI	00	00
7	6	5	4	3	2	1	0

2-way Set Associative Cache						
Index	Tag	Valid	Bytes(0-3)	Tag	Valid	Bytes(0-3)
0	0	0	00 00 00 00	0	0	00 00 00 00
1	(3)	(1)	02 B3 (8F) B6	<del>8</del> A	<del>1</del> 1	00 00 00 00 4D A2
2	<del>8</del> A	<del>1</del> 1	00 00 00 00 5D	(8)	(1)	EF 09 (86) 2A
3	(B)	(1)	C5 EC 76 4E	<del>3</del> 3	<del>1</del> 1	00 00 00 00 19

Address Reference Sequence						
Addr	0x87	0x88	0xA5	0xA8	0x36	0x3E
Hit?	N	Y	N	N	Y	N

Memory Contents																
Address	Bytes x0 to xF relative to address															
	x0	x1	x2	<u>x3</u>	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x00	AA	03	3E	CD	38	16	7B	ED	5A	8E	4C	DF	18	FB	B7	12
0x10	A9	76	2B	EE	BC	91	D5	92	80	BA	9B	F6	48	16	81	0A
0x20	75	F7	3F	C6	9E	3A	0F	DA	00	4C	B6	A8	92	04	E5	2E
<u>0x30</u>	E0	22	<u>19</u>	3A	02	B3	8F	B6	25	31	E1	02	18	09	73	02
0x40	86	B8	F0	C6	AA	29	AE	16	76	46	80	6E	13	EA	A8	66
0x50	04	2A	32	6A	B1	86	56	0E	96	30	47	F2	F8	1D	42	30
0x60	2F	7E	3D	A8	27	95	A4	74	07	11	6B	D8	C7	B7	AF	C2
0x70	D6	A4	89	92	FD	FE	D6	DA	DE	D5	CD	4A	7C	68	3A	1A
<u>0x80</u>	ED	32	0A	<u>A2</u>	BF	80	1D	FC	EF	09	86	2A	25	44	6F	1A
0x90	1E	C2	AE	60	5C	3E	DF	F2	25	CF	84	DA	F1	6B	DC	DE
<u>0xA0</u>	<u>5D</u>	<u>4D</u>	F7	DA	69	C2	8C	74	A8	CE	7F	DA	FA	93	EB	48
0xB0	61	C6	5E	74	03	97	BA	62	F8	11	72	12	C5	EC	76	4E
0xC0	17	52	75	2C	62	89	EF	18	BB	7D	8C	7C	26	57	7F	C2
0xD0	54	9E	1E	FA	DC	81	B2	14	B6	1F	7B	44	10	F5	B8	2E
0xE0	14	9A	0D	4A	C8	1D	E6	6E	F3	38	F3	5C	6C	8F	BD	A8
0xF0	32	21	1C	2C	22	C2	DC	34	BA	DD	37	D8	E7	A2	39	BA

Add. 0x87:

7	6	5	4	3	2	1	0
1	0	0	0	0	1	1	1

- CO: 0x3 (11)
- CI: 0x1 (01)
- CT: 0x8 (1000)
- Cache Hit: N

Add 0x88:

7	6	5	4	3	2	1	0
1	0	0	0	1	0	1	0

- CO: 0x2 (10)
- CI: 0x2 (10)
- CT: 0x8 (1000)
- Cache Hit: Y
- Cache byte returned: 0x86

Add. 0xA5:

7	6	5	4	3	2	1	0
1	0	1	0	0	1	0	1

- CO: 0x1 (01)
- CI: 0x1 (01)
- CT: 0xA (1010)
- Cache Hit: N

Add. 0xA8:

7	6	5	4	3	2	1	0
1	0	1	0	1	0	0	0

- CO: 0x0 (00)
- CI: 0x2 (10)
- CT: 0xA (1010)
- Cache Hit: N

Add. 0x36:

7	6	5	4	3	2	1	0
0	0	1	1	0	1	1	0

- CO: 0x2 (10)
- CI: 0x1 (01)
- CT: 0x3 (0011)
- Cache Hit: Y
- Cache Byte returned: 0x8F

Add. 0x3E:

7	6	5	4	3	2	1	0
0	0	1	1	1	1	1	0

- CO: 0x2 (10)
- CI: 0x3 (11)
- CT: 0x3 (0011)
- Cache Hit: N